

Lab 7 - Kubernetes

Initialize cluster

1. To free up resources, stop all vagrant VMs in devop-lab environment.

- Clone the following Git repository :

<https://github.com/brahimhamdi/k8s-lab>

- In k8s-lab directory, execute following command to deploy k8s vagrants VMs :

vagrant up

2. Kubernetes is already installed on all vagrant VMs. On master VM, initialize the cluster.

sudo kubeadm init --apiserver-advertise-address=192.168.56.10 --pod-network-cidr=10.244.0.0/16

```
vagrant@kubernetes-control-plane:~$ sudo kubeadm init --apiserver-advertise-address=192.168.56.10 --pod-network-cidr=10.244.0.0/16
[init] Using Kubernetes version: v1.27.4
[preflight] Running pre-flight checks
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action in beforehand using 'kubeadm config images pull'
```

- If no errors, what's the output of the initializing command ?

```
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy
```

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, if you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 192.168.56.10:6443 --token d8qczx.mki4o33z0k11i0ih \
--discovery-token-ca-cert-hash sha256:15018c12c56b176cd793c7e16d6f2df85372427409f4d35f218b4c118cef8149
vagrant@kubernetes-control-plane:~$
```

- To start using your cluster as regular user, apply next commands

```
vagrant@k8s-master:~$ mkdir -p $HOME/.kube
vagrant@k8s-master:~$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
vagrant@k8s-master:~$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
vagrant@k8s-master:~$
```

- Apply flannel yaml file.

kubectl apply -f <https://raw.githubusercontent.com/flannel-io/flannel/master/Documentation/kube-flannel.yml>

```
vagrant@k8s-master:~$ kubectl apply -f https://raw.githubusercontent.com/flannel-io/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
vagrant@k8s-master:~$
```

3. Check the cluster info.

```
vagrant@k8s-master:~$ kubectl cluster-info
Kubernetes control plane is running at https://192.168.56.10:6443
CoreDNS is running at https://192.168.56.10:6443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy
```

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.

```
vagrant@k8s-master:~$
```

- How kubernetes components looks like ?

```
vagrant@k8s-master:~$ kubectl get namespaces
NAME                STATUS    AGE
default              Active    13m
kube-flannel         Active    6m16s
kube-node-lease      Active    13m
kube-public          Active    13m
kube-system          Active    13m
vagrant@k8s-master:~$ kubectl get pod -n kube-system
NAME                                READY   STATUS    RESTARTS   AGE
coredns-5d78c9869d-brhv4           0/1     Pending   0           13m
coredns-5d78c9869d-nnhh8           0/1     Pending   0           13m
etcd-k8s-master                    1/1     Running   10          13m
kube-apiserver-k8s-master           1/1     Running   19          13m
kube-controller-manager-k8s-master 1/1     Running   20          13m
kube-proxy-54pxj                   1/1     Running   0           13m
kube-scheduler-k8s-master           1/1     Running   20          13m
vagrant@k8s-master:~$
```

- What is the IP address of DNS systems ?

```
vagrant@k8s-master:~$ kubectl get pod -n kube-system -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE                NOMINATED NODE
coredns-5d78c9869d-5gxxx           1/1     Running   0           19m   10.244.0.3      kube-control-plane   <none>
coredns-5d78c9869d-lfdrf           1/1     Running   0           19m   10.244.0.2      kube-control-plane   <none>
etcd-kube-control-plane             1/1     Running   0           20m   192.168.56.10   kube-control-plane   <none>
```

4. Join all nodes to the cluster.

- On the master, check that all nodes are ready on the cluster.

sudo kubeadm token create --print-join-command

```
vagrant@k8s-control-plane:~$ sudo kubeadm token create --print-join-command
kubeadm join 192.168.56.10:6443 --token 0n48cr.x0tzlke0pt4un26e --discovery-token-ca-cert-hash sha256:15018c12c56b176cd793c7e16d6f2df85372427409f4d35f218b4c118cef8149
vagrant@k8s-control-plane:~$
```

```
vagrant@k8s-control-plane:~$ exit
logout
brahim@Training:~/k8s-lab$ vagrant ssh kube-node2
vagrant@k8s-node2:~$ sudo kubeadm join 192.168.56.10:6443 --token 0n48cr.x0tzlke0pt4un26e --discovery-token-ca-cert-hash sha256:15018c12c56b176cd793c7e16d6f2df85372427409f4d35f218b4c118cef8149
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...
```

This node has joined the cluster:
 * Certificate signing request was sent to apiserver and a response was received.
 * The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

```
vagrant@k8s-node2:~$ exit
logout
brahim@Training:~/k8s-lab$ vagrant ssh kube-control-plane
Last login: Tue Aug 8 12:42:42 2023 from 10.0.2.2
vagrant@k8s-control-plane:~$ kubectl get node -o wide
NAME                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
k8s-control-plane   Ready     control-plane   29m   v1.27.4   192.168.56.10   <none>        Ubuntu 20.04.5 LTS   5.4.0-139-generic   containerd
k8s-node1           Ready     <none>        2m6s   v1.27.4   192.168.56.11   <none>        Ubuntu 20.04.5 LTS   5.4.0-139-generic   containerd
k8s-node2           NotReady  <none>        13s    v1.27.4   192.168.56.12   <none>        Ubuntu 20.04.5 LTS   5.4.0-139-generic   containerd
vagrant@k8s-control-plane:~$
vagrant@k8s-control-plane:~$
```

```
vagrant@k8s-control-plane:~$ sudo kubeadm token create --print-join-command
kubeadm join 192.168.56.10:6443 --token 0n48cr.x0tzlke0pt4un26e --discovery-token-ca-cert-hash sha256:15018c12c56b176cd793c7e16d6f2df85372427409f4d35f218b4c118cef8149
vagrant@k8s-control-plane:~$
vagrant@k8s-control-plane:~$ exit
logout
```

```
brahim@Training:~/k8s-lab$ vagrant ssh kube-node1
vagrant@k8s-node1:~$ sudo kubeadm join 192.168.56.10:6443 --token 0n48cr.x0tzlke0pt4un26e --discovery-token-ca-cert-hash sha256:15018c12c56b176cd793c7e16d6f2df85372427409f4d35f218b4c118cef8149
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...
```

This node has joined the cluster:
 * Certificate signing request was sent to apiserver and a response was received.
 * The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

```
vagrant@k8s-node1:~$
vagrant@k8s-node1:~$ exit
logout
brahim@Training:~/k8s-lab$ vagrant ssh kube-control-plane
Last login: Tue Aug 8 12:09:11 2023 from 10.0.2.2
vagrant@k8s-control-plane:~$ kubectl get node -o wide
NAME                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
k8s-control-plane   Ready     control-plane   27m   v1.27.4   192.168.56.10   <none>        Ubuntu 20.04.5 LTS   5.4.0-139-generic   containerd
k8s-node1           NotReady  <none>        48s    v1.27.4   192.168.56.11   <none>        Ubuntu 20.04.5 LTS   5.4.0-139-generic   containerd
vagrant@k8s-control-plane:~$
```

Manage pods

5. Create a yaml file for a *hasher* pod.

```
vagrant@k8s-master:~$ vim hasher.yaml
vagrant@k8s-master:~$ cat hasher.yaml
apiVersion: v1
kind: Pod
metadata:
  name: hasher
spec:
  containers:
  - name: hasher
    image: brahimhamdi/hasher
vagrant@k8s-master:~$ kubectl apply -f hasher.yaml
pod/hasher created
vagrant@k8s-master:~$
vagrant@k8s-master:~$ kubectl get pod -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP        NODE      NOMINATED NODE   READINESS GATES
hasher    0/1     ContainerCreating   0       6s     <none>    k8s-worker2   <none>           <none>
```

6. Apply the yaml file.

- On which node the pod is created ?
- What is the pod's IP address ?
- What is the container's name and ID ?
- What is the image's name and ID ?

```
vagrant@k8s-master:~$ kubectl describe pod/hasher
Name:          hasher
Namespace:     default
Priority:       0
Service Account: default
Node:          k8s-worker2/192.168.205.102
Start Time:    Tue, 13 Jun 2023 13:36:56 +0000
Labels:        <none>
Annotations:   <none>
Status:        Running
IP:            10.244.2.63
IPs:
  IP: 10.244.2.63
Containers:
  hasher:
    Container ID:  containerd://755099de8a6f3fd62d23b228d212c78f755a0ecb6a867c755776fba923f84c87
    Image:         brahimhamdi/hasher
    Image ID:      docker.io/brahimhamdi/hasher@sha256:a37377f07840109415eb7df07ae830bc617d0f3ac3c98c904b7a8647868785f5
    Port:         <none>
    Host Port:    <none>
    State:        Running
      Started:    Tue, 13 Jun 2023 13:38:38 +0000
    Ready:        True
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-9fkd1 (ro)
Conditions:
  Type             Status
  Initialized       True
  Ready            True
  ContainersReady  True
  PodScheduled     True
Volumes:
```

7. Remove the pod from the cluster.

```
vagrant@k8s-master:~$ kubectl delete pod hasher
pod "hasher" deleted
vagrant@k8s-master:~$ kubectl get pod
No resources found in default namespace.
vagrant@k8s-master:~$
```

Manage deployments and services

8. Create yaml file to describe *dockercoins* application deployment.

```
apiVersion: v1
kind: Namespace
metadata:
  name: dockercoins
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: worker
  namespace: dockercoins
spec:
  # replicas: 3
  selector:
    matchLabels:
      app: dockercoins
      tier: worker
  template:
    metadata:
      labels:
        app: dockercoins
        tier: worker
    spec:
      containers:
        - name: worker
          image: brahimhamdi/worker
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: rng
  namespace: dockercoins
spec:
  # replicas: 3
  selector:
    matchLabels:
      app: dockercoins
      tier: rng
  template:
    metadata:
      labels:
        app: dockercoins
        tier: rng
    spec:
      containers:
        - name: rng
          image: brahimhamdi/rng
---
apiVersion: apps/v1
kind: Deployment
metadata:
```

```
name: hasher
namespace: dockercoins
spec:
# replicas: 3
selector:
  matchLabels:
    app: dockercoins
    tier: hasher
template:
  metadata:
    labels:
      app: dockercoins
      tier: hasher
  spec:
    containers:
      - name: hasher
        image: brahimhamdi/hasher
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: redis
  namespace: dockercoins
spec:
# replicas: 3
selector:
  matchLabels:
    app: dockercoins
    tier: redis
template:
  metadata:
    labels:
      app: dockercoins
      tier: redis
  spec:
    containers:
      - name: redis
        image: redis
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: webui
  namespace: dockercoins
spec:
# replicas: 3
selector:
  matchLabels:
    app: dockercoins
    tier: webui
template:
  metadata:
    labels:
      app: dockercoins
      tier: webui
  spec:
    containers:
      - name: webui
```

```
    image: brahimhamdi/webui
```

```
---
```

```
apiVersion: v1
kind: Service
metadata:
  name: rng
  namespace: dockercoins
spec:
  selector:
    app: dockercoins
    tier: rng
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: ClusterIP
```

```
---
```

```
apiVersion: v1
kind: Service
metadata:
  name: hasher
  namespace: dockercoins
spec:
  selector:
    app: dockercoins
    tier: hasher
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: ClusterIP
```

```
---
```

```
apiVersion: v1
kind: Service
metadata:
  name: redis
  namespace: dockercoins
spec:
  selector:
    app: dockercoins
    tier: redis
  ports:
    - protocol: TCP
      port: 6379
      targetPort: 6379
  type: ClusterIP
```

```
---
```

```
apiVersion: v1
kind: Service
metadata:
  name: webui
  namespace: dockercoins
spec:
  selector:
    app: dockercoins
    tier: webui
  ports:
    - protocol: TCP
```

port: 80
 targetPort: 80
 nodePort: 30001
 type: NodePort

9. Apply the yaml file and check the application.

```
vagrant@k8s-master:~$ kubectl apply -f dockercoins.yaml
namespace/dockercoins created
deployment.apps/worker created
deployment.apps/rng created
deployment.apps/hasher created
deployment.apps/redis created
deployment.apps/webui created
service/rng created
service/hasher created
service/redis created
service/webui created
vagrant@k8s-master:~$
```

```
vagrant@k8s-master:~$ kubectl get all -n dockercoins
```

NAME	READY	STATUS	RESTARTS	AGE
pod/hasher-7f9d944db9-d2xbl	1/1	Running	0	3m24s
pod/redis-78579d7b98-l4sp2	1/1	Running	0	3m24s
pod/rng-544477487c-c8dn8	1/1	Running	0	3m24s
pod/webui-c9697458-8857v	1/1	Running	0	3m24s
pod/worker-5f7877988-frxhg	1/1	Running	0	3m24s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/hasher	ClusterIP	10.104.236.71	<none>	80/TCP	3m23s
service/redis	ClusterIP	10.100.243.140	<none>	6379/TCP	3m23s
service/rng	ClusterIP	10.111.114.165	<none>	80/TCP	3m23s
service/webui	NodePort	10.108.202.27	<none>	80:30001/TCP	3m23s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/hasher	1/1	1	1	3m24s
deployment.apps/redis	1/1	1	1	3m24s
deployment.apps/rng	1/1	1	1	3m24s
deployment.apps/webui	1/1	1	1	3m24s
deployment.apps/worker	1/1	1	1	3m24s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/hasher-7f9d944db9	1	1	1	3m24s
replicaset.apps/redis-78579d7b98	1	1	1	3m24s
replicaset.apps/rng-544477487c	1	1	1	3m24s
replicaset.apps/webui-c9697458	1	1	1	3m24s
replicaset.apps/worker-5f7877988	1	1	1	3m24s


```
vagrant@k8s-master:~$
```



```
vagrant@k8s-master:~$ vim dockercoins.yaml
vagrant@k8s-master:~$ kubectl apply -f dockercoins.yaml
namespace/dockercoins unchanged
deployment.apps/worker configured
deployment.apps/rng unchanged
deployment.apps/hasher unchanged
deployment.apps/redis unchanged
deployment.apps/webui unchanged
service/rng unchanged
service/hasher unchanged
service/redis unchanged
service/webui unchanged
vagrant@k8s-master:~$
```

← → ↻ 🏠 192.168.205.101:30001/index.html

DockerCoin Miner WebUI



Current mining speed: ~28.1 hashes/second ([Tweet this!](#))

